

INTRODUCTION

In this project you will implement a multiprocessor operating system simulator using a popular user space threading library for Linux called **pthread**s. The framework for the multi threaded OS simulator is nearly complete, but missing one critical component: the CPU scheduler! Your task is to implement the CPU scheduler, using three different scheduling algorithms.

IMPORTANT: PLEASE DO NOT SHARE SOURCE CODE WITH OTHER STUDENTS. DO NOT ATTEMPT TO USE SOLUTIONS YOU FIND ON THE INTERNET, OR COPY SOLUTIONS FROM THOSE WHO TOOK THE CLASS IN PRIOR SEMESTERS. WE ARE USING A SOPHISTICATED SOFTWARE TO CATCH ALL SUCH CASES, AND FILLING AN ACADEMIC MIS-CONDUCT IS NOT FUN FOR ANYONE.

NOTE: MAKE SURE THAT MULTIPLE CPU CORES ARE TURNED ON IN YOUR VIRTUAL MACHINE

We have provided you with source files that constitute the framework for your simulator. You will only need to modify **answers.txt** and **student.c**. However, just because you are only modifying two files doesn't mean that you should ignore the other ones - there is helpful information in the other files. Information about using the pthreads library is given in **Problem 0**. We have provided you these files:

- Makefile - Working one provided to you; **Modify at your own peril**
- os-sim.c - Code for the operating system simulator which calls your CPU scheduler
- os-sim.h - Header file for the simulator which calls your CPU scheduler
- process.c - Descriptions of the simulated processes
- process.h - Header file for the process data
- student.h - Header file for your code to interface with the OS scheduler
- student.c - This file contains stub functions for your CPU scheduler. **You will be modifying this file**
- answers.txt - A blank file where you will be answering some questions

Scheduling Algorithms

For your simulator, you will implement the following scheduling algorithms:

- *First In, First Out (FIFO)* - Runnable processes are kept in a ready queue. FIFO is non-preemptive; once a process begins running on a CPU, it will continue until it either completes or blocks for I/O
- *Round Robin* - Similar to FIFO, except preemptive. Each process is assigned a time slice when it is scheduled. At the end of the time slice, if the process is still running, the process is preempted, and moved to the tail of the ready queue.
- *Static Priority* - The processes with the highest priorities always get the CPU. Lower-priority processes may be preempted if a process with higher priority becomes runnable.

Process States

In our OS simulation, there are five possible states for a process, which are described in the `process_state_t` enum in the `os-sim.h`:

- NEW - The process is being created, and has not yet begun executing
- READY - The process is ready to execute, and is waiting to be scheduled on a CPU

- **RUNNING** - The process is currently executing on a CPU
- **WAITING** - The process has temporarily stopped executing, and is waiting on an I/O request to complete
- **TERMINATED** - The process has completed

There is a field name `state` in the PCB, which must be updated with the current state of the process. The simulator will use this field to collect statistics.

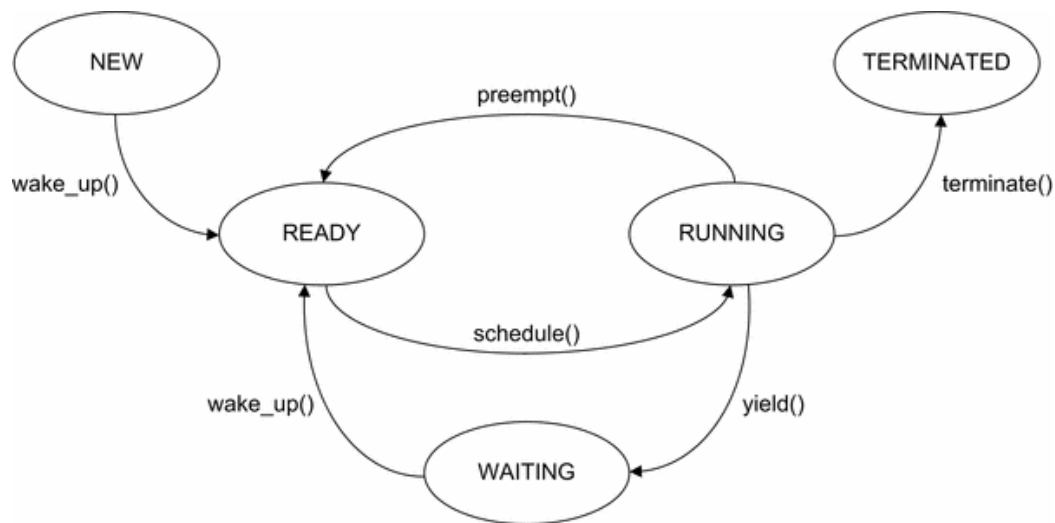


Figure 1: Process States

The Ready Queue

On most systems there are a large number of processes, but only one or two CPUs on which to execute them. When there are more processes ready to execute than CPUs, processes must wait in the READY state until a CPU becomes available. To keep track of the processes waiting to execute, we keep a ready queue of the processes in the READY state.

Since the ready queue is accessed by multiple processors, which may add and remove processes from it, the ready queue must be protected by some form of synchronization—for this project, you will use a mutex lock.

Scheduling Processes

`schedule()` is the core function of the CPU scheduler. It is invoked whenever a CPU becomes available for running a process. `schedule()` must search the ready queue, select a runnable process, and call the `context_switch()` function to switch the process onto the CPU.

There is a special process, the idle process, which is scheduled whenever there are no processes in the READY state.

CPU Scheduler Invocation

There are four events which will cause the simulator to invoke `schedule()`:

1. `yield()` - A process completes its CPU operations and yields the processor to perform an I/O request

2. `wake_up()` - A process that previously yielded completes its I/O request, and is ready to perform CPU operations. `wake_up()` is also called when a process in the NEW state becomes runnable
3. `preempt()` - When using a Round-Robin or Static Priority scheduling algorithms, a CPU bound process may be preempted before it completes its CPU operations
4. `terminate()` - A process exits or is killed

The CPU scheduler also contains one other important function: `idle()`; This functions simulates the idle process. In the real world, the idle process puts the processor in a low power mode and waits. For our OS simulation, you will make use of the **pthread condition** variable to block the thread until a process enters the ready queue again.

The Simulator

We will use pthreads to simulate an operating system on a multiprocessor computer. We will use one thread per CPU and one thread as a “supervisor” for our simulation. The CPU threads will simulate the currently-running processes on each CPU, and the supervisor thread will print output and dispatch events to the CPU threads.

Since the code you write will be called from multiple threads, the CPU scheduler you write must be thread-safe! This means that all data structures you use, including your ready queue, must be protected using mutexes.

The number of CPUs is specified as a command-line parameter to the simulator. For this project, you will be performing experiments with 1, 2, and 4 CPU simulations.

Also, for demonstration purposes, the simulator executes much slower than a real system would. In the real world, a CPU burst might range from one to a few hundred milliseconds, whereas in this simulator, they range from 0.2 to 2.0 seconds.

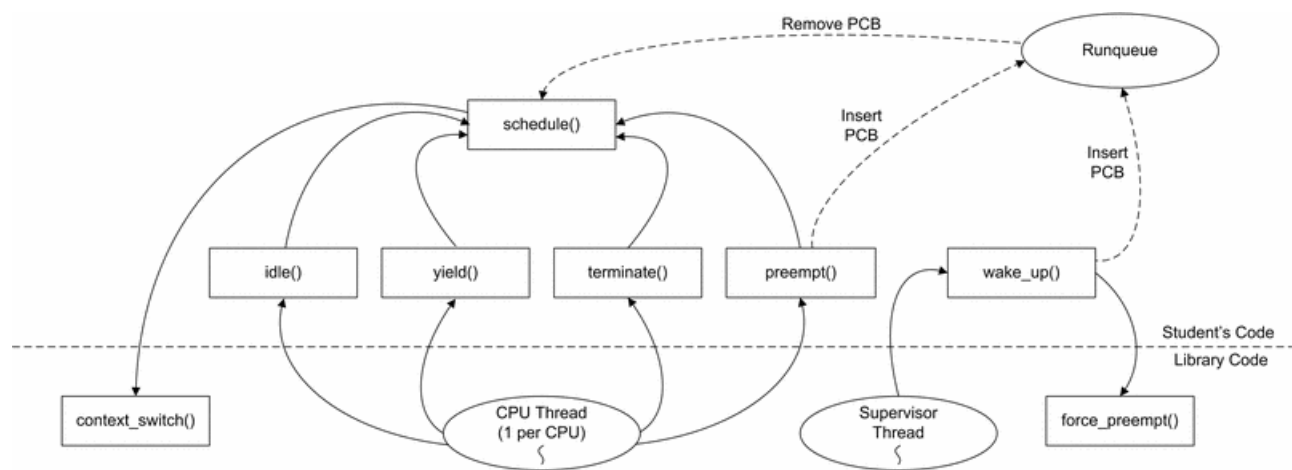


Figure 2: Simulator Function Calls

Sample Output

Compile and run the simulator with `./os-sim 2`. After a few seconds, hit Control-C to exit. You will see the output below:

Time =====	Ru ==	Re ==	Wa ==	CPU0 =====	CPU1 =====	I/O Queue =====
0.0	0	0	0	(IDLE)	(IDLE)	<<
0.1	0	0	0	(IDLE)	(IDLE)	<<
0.2	0	0	0	(IDLE)	(IDLE)	<<
0.3	0	0	0	(IDLE)	(IDLE)	<<
0.4	0	0	0	(IDLE)	(IDLE)	<<
0.5	0	0	0	(IDLE)	(IDLE)	<<
0.6	0	0	0	(IDLE)	(IDLE)	<<
0.7	0	0	0	(IDLE)	(IDLE)	<<
0.8	0	0	0	(IDLE)	(IDLE)	<<
0.9	0	0	0	(IDLE)	(IDLE)	<<
1.0	0	0	0	(IDLE)	(IDLE)	<<
.....						

The simulator generates a Gantt Chart, showing the current state of the OS at every 100ms interval. The leftmost column shows the current time, in seconds. The next three columns show the number of Running, Ready, and Waiting processes, respectively. The next two columns show the process currently running on each CPU. The rightmost column shows the processes which are currently in the I/O queue, with the head of the queue on the left and the tail of the queue on the right.

As you can see, nothing is executing. This is because we have no CPU scheduler to select processes to execute! Once you complete Problem 1 and implement a basic FIFO scheduler, you will see the processes executing on the CPUs.

Test Processes

For this simulation, we will use a series of eight test processes, five CPU-bound and three I/O bound. For simplicity, we have labeled each process starting with a "C" or "I" to indicate CPU or I/O bound respectively. The table below shows a detailed breakdown of the processes.

For this project, priorities range from 0 to 10, with 10 being the highest priority. Note that the I/O-bound processes have been given higher priorities than the CPU-bound processes.

Table 2: Process Descriptions

PID	Process Name	CPU or I/O bound	Priority	Start Time
0	lapache	I/O-bound	8	0.0 s
1	lbash	I/O-bound	7	1.0 s
2	Imozilla	I/O-bound	7	2.0 s
3	Ccpu	CPU-bound	5	3.0 s
4	Cgcc	CPU-bound	1	4.0 s
5	Cspice	CPU-bound	2	5.0 s
6	Cmysql	CPU-bound	3	6.0 s
7	Csim	CPU-bound	4	7.0 s

Problem 0: pthreads Review

[0 points]

Spend some time and take a look at the pthreads documentation. Make a small multi-threaded program where two threads print the numbers between 1 and 1000. This will help you better understand the life cycle of threads.

You should also use the below resources to get a better idea of pthreads:

- man pages for all the relevant pthread library calls. In particular, look at `pthread_mutex_init`, `pthread_mutex_lock()`, `pthread_cond_init`, `pthread_cond_broadcast`, and `pthread_cond_wait`.
- <http://www.llnl.gov/computing/tutorials/pthreads/>

Note: When you get to using `pthread_cond_wait()`, use a `while` loop instead of an `if` statement to enclose the call to the function. If you look carefully, the pthread documentation says that `pthread_cond_wait` function may return without having acquired the lock. The `while` makes sure that the condition is checked before continuing with the execution, ensuring that we acquire the lock. Using an `if` may cause completely untraceable bugs in your programs.

Problem 1: FIFO Scheduler

A. [50 points]

Implement the CPU scheduler using the FIFO algorithm. You may do this however you like, however, we suggest the following:

- Implement a thread-safe ready queue using a linked list. A linked list implementation will allow you to reuse this ready queue for the Round-Robin and Static Priority scheduling algorithms. Make sure that the mutex variable you use is different from the mutex variable used for the CPU's (i.e. It is not the same as the current mutex already defined in the framework code). (**Hint:** Look at the `pcb_t` struct in `os-sim.h`.)
- Implement the `yield()`, `wake_up()`, and `terminate()` functions in `student.c`. `preempt()` is not necessary for this stage of the project. See the introduction and comments in the code for the proper behavior of these events.
- Implement `idle()`. This function must wait on a condition variable that is signaled whenever a process is added to the ready queue.

- Implement `schedule()`. This function should extract the first process in the ready queue, then call `context_switch()` to select the process to execute. If there are no runnable processes, `schedule()` should call `context_switch` with a `NULL` pointer as parameter to execute the idle process.

Before you begin writing code, look at the contents of the file `os-sim.h` for a list of function prototypes and descriptions of the currently used data structures.

Once you successfully complete this portion of the project, make and test your code with `./os-sim 1`. You should see an output similar to the following:

Time =====	Ru ==	Re ==	Wa ==	CPU0 =====	I/O Queue =====
0.0	0	0	0	(IDLE)	<<
0.1	1	0	0	Iapache	<<
0.2	1	0	0	Iapache	<<
0.3	1	0	0	Iapache	<<
0.4	0	0	1	(IDLE)	<Iapache <
0.5	0	0	1	(IDLE)	<Iapache <
0.6	1	0	0	Iapache	<<
0.7	1	0	0	Iapache	<<
0.8	1	0	0	Iapache	<<
0.9	1	0	0	Iapache	<<
1.0	0	0	1	(IDLE)	<Iapache <
1.1	1	0	1	Ibash	<Iapache <
1.2	1	0	1	Ibash	<Iapache <
1.3	1	0	1	Ibash	<Iapache <
1.4	1	0	1	Ibash	<Iapache <
1.5	1	0	1	Iapache	<Ibash <
1.6	1	0	1	Iapache	<Ibash <
1.7	0	0	2	(IDLE)	<Ibash Iapache <
1.8	0	0	2	(IDLE)	<Ibash Iapache <
1.9	0	0	2	(IDLE)	<Ibash Iapache <
2.0	0	0	0	Ibash	<Iapache <

```

.....
66.9    1    1    0        Ibash        <<
67.0    1    1    0        Ibash        <<
67.1    1    1    0        Ibash        <<
67.2    1    0    0        Imozilla     <<
67.3    1    0    0        Imozilla     <<
67.4    1    0    0        Imozilla     <<
67.5    1    0    0        Imozilla     <<

```

```

# of Context Switches: 97
Total execution time: 67.6 s
Total time spent in READY state: 389.9 s
(These numbers may be slightly different for you)

```

Important Information:

- Be sure to update the state field of the PCB. The library will read this field to generate the Running, Ready, and Waiting columns of the output. These are also used to generate the statistics and the end of the simulation.
- Four of the five entry points in the scheduler (`idle()`, `yield()`, `terminate()`, and `preempt()`) should cause a new process to be scheduled on the CPU. In your handlers, be sure to call `schedule()`, which will select a runnable process, and then call `context_switch()`. When these four functions return, the library will simulate the process selected by `context_switch()`.
- `context_switch()` takes a time slice parameter, which is used for preemptive scheduling algorithms. Since FIFO is non-preemptive, use -1 for this parameter to give the process an infinite time slice.

B. [10 points]

Run your OS simulation with 1, 2, and 4 CPUs. Compare the total execution time of each. Is there a linear relationship between the number of CPUs and total execution time? Why or why not?

Problem 2: Round-Robin Scheduler

A. [10 points]

Add Round-Robin scheduling functionality to your code. You should modify `main()` to add a command line option, `-r`, which selects the Round-Robin scheduling algorithm, and accepts a parameter, the length of the time slice. For this project, time slices are measured in tenths of seconds. E.g.:

```
./os-sim <# of CPUs> -r 5
```

Should run a Round-Robin scheduler with time slices of 500 ms. While:

```
./os-sim <# of CPUs>
```

Should continue to run the FIFO scheduler.

Make sure that you also implement the `preempt()` function. To specify a time slice when scheduling a process, use the time slice parameter of `context_switch()`. The simulator will automatically preempt the process and call your `preempt()` handler when a process finishes executing on the CPU for the length of the time slice without terminating or yielding for I/O.

B. [10 points]

Run your Round-Robin scheduler with time slices of 800ms, 600ms, 400ms, and 200ms. Use only one CPU for your tests. Compare the statistics at the end of the simulation. You will see that the total waiting time decreases with shorter time slices. However, in a real OS, the shortest time slice may not be the best choice. Explain why that is the case?

Static Priority Scheduler

A. [10 points]

Add Static-Priority scheduling support to your code. Modify `main()` to accept the “-p” parameter to select the Static Priority algorithm. The “-r” and default FIFO scheduler should continue to work.

The scheduler should use the priority specified in the `static_priority` field of the PCB. This priority is a value from 0 to 10, with 0 being the lowest priority and 10 being the highest priority.

For Static Priority scheduling, you will need to make use of the `current[]` array and `force_preempt()` function. The `current[]` array should be used to keep track of the process currently executing on each CPU. Since this array is accessed by multiple CPU threads, it must be protected by a mutex. `current_mutex` has been provided to you for this purpose.

The `force_preempt()` function preempts a running process before its time slice expires. Your `wake_up()` handler should make use of this function to preempt a lower priority process when a higher priority process needs a CPU.

B. [10 points]

The Shortest-Job First (SJF) scheduling algorithm is proven to have the optimal average waiting time. However, it is not feasible to implement in a typical scheduler, since the scheduler does not have advanced knowledge of the length of each CPU burst.

Run each of your three scheduling algorithms (using one CPU), and compare the total waiting times. Which algorithm is the closest approximation of SJF? Why?

Assignment Submission

Note: Each problem has two parts (labeled A and B). The first is the actual implementation, and the second is a conceptual question that is to be answered after running some tests. Make sure you complete both.

We have provided you with a `make submit` command in the Makefile. Use it to generate a tar ball that can be submitted.

Please Note:

- Make sure that the tar ball contains the following item:
 - answers.txt - Short answers for part B of all the problems.
 - Makefile - Working one has been provided to you;
 - os-sim.c - Code for the operating system simulator
 - os-sim.h - Header file for the simulator
 - process.c - Descriptions of the simulated processes
 - process.h - Header file for the process data
 - student.c - Your code for the scheduler
 - student.h - Header file for your scheduler code

We suggest untarring the tar ball to make sure that the above contents are all present

- **If you code does not compile, you will get a zero!**
- Keep your answers for part B of all the problems detailed enough to cover the question, including support from simulator results if appropriate. Don't write a book; however when in doubt, err on the side of giving us too much information.